Instructor: Mithun Prasad, PhD
Email: miprasad@Microsoft.com

Microsoft

WHAT YOU WILL BE ABLE TO DO AFTER THIS TRAINING

Build a Data Science experiment using ML studio.

Gain familiarity with Data Science components of the studio.

Customize Data Science components in the studio.

Microsoft

MACHINE LEARNING 101

The first way of thinking about ML is by the type of information or input given to a system.

1. **Supervised learning** – we get the data and the labels e.g. linear regression
2. **Unsupervised learning** – only get the data (no labels) e.g. clustering
3. **Reinforcement learning** – reward/penalty based information (feedback)

Another way of categorizing ML approaches, is to the desired output:

1. **Classification** (e.g. decision tree)
2. **Regression** (e.g. linear regression)
3. **Clustering** (e.g. k-means)
4. **Density estimation** (e.g. histograms)
5. **Dimensionality reduction** (e.g. principal component analysis

MACHINE LEARNING CAPABILITIES

Which category
(Classification)

How much/many
(Regression)

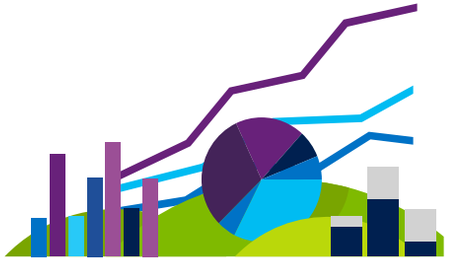

Which group
(Clustering, Recommender)

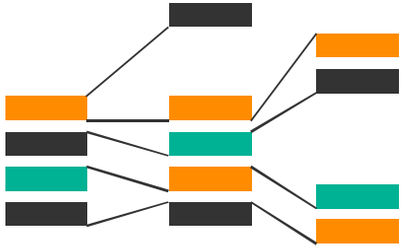

Is it odd
(Anomaly)

Which action
(Reinforcement Learning)

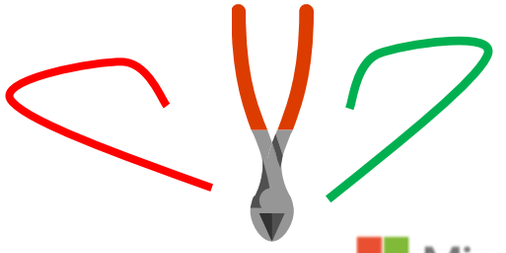

Microsoft

MACHINE LEARNING 101

Term	Definition
Training set	set of data used to learn a model
Test set	set of data used to test a model
Feature	a variable (continuous, discrete, categorical, etc.) aka column
Target	Label (associated with dependent variable, what we predict)
Learner	Model or algorithm
Fit, Train	Learn a model with an ML algorithm using a training set
Predict	w/ supervised learning, give a label to an unknown datum(data). w/unsupervised decide if new data is weird, in which group, or what to do next with the new data
Accuracy	percentage of correct predictions ((TP + TN)/ total)
Precision	Percentage of correct positive predictions (TP/ (FP + TP))
Recall	Percentage of positive cases caught (TP/ (FN + TP))

CORTANA INTELIGENCE IN A SENTENCE

Cortana Intelligence is a Platform and a Process to perform advanced analytics from start to finish

Microsoft

THE CORTANA INTELLIGENCE PLATFORM

Cortana, Cognitive Services, Bot Framework

Power BI

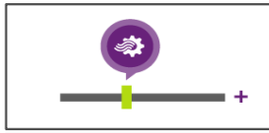

Azure Stream Analytics

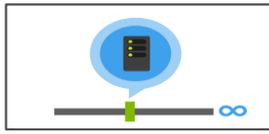

Azure HDInsight

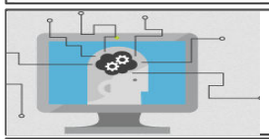

Azure Machine Learning and MRS

Azure SQL DB, Data Warehouse, DocumentDB

Azure Data Lake

Azure Event Hubs

Azure Data Factory

Azure Data Catalog

Microsoft Azure

Microsoft

AZURE ML ELEMENTS

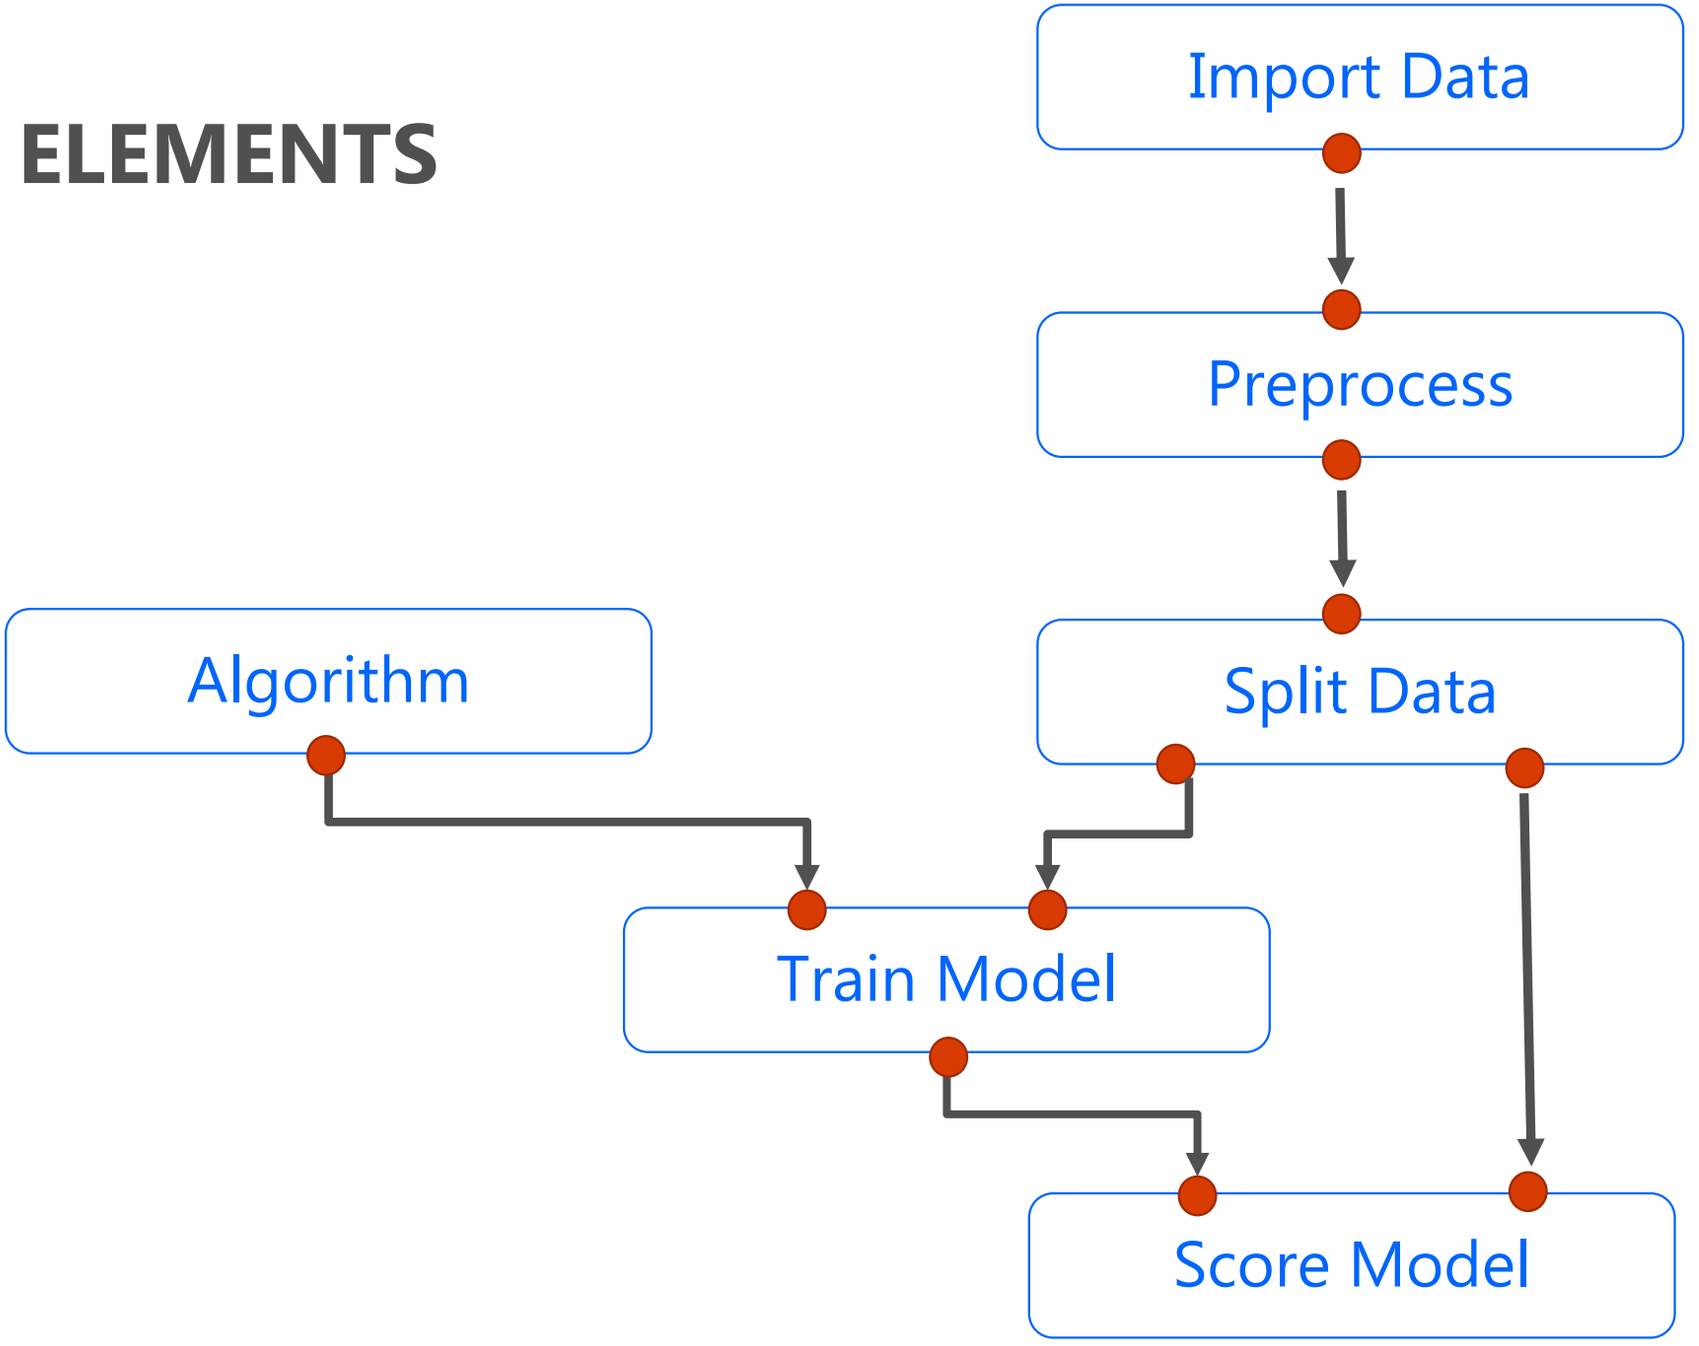

AZURE ML STUDIO AND THE TEAM DATA SCIENCE PROCESS

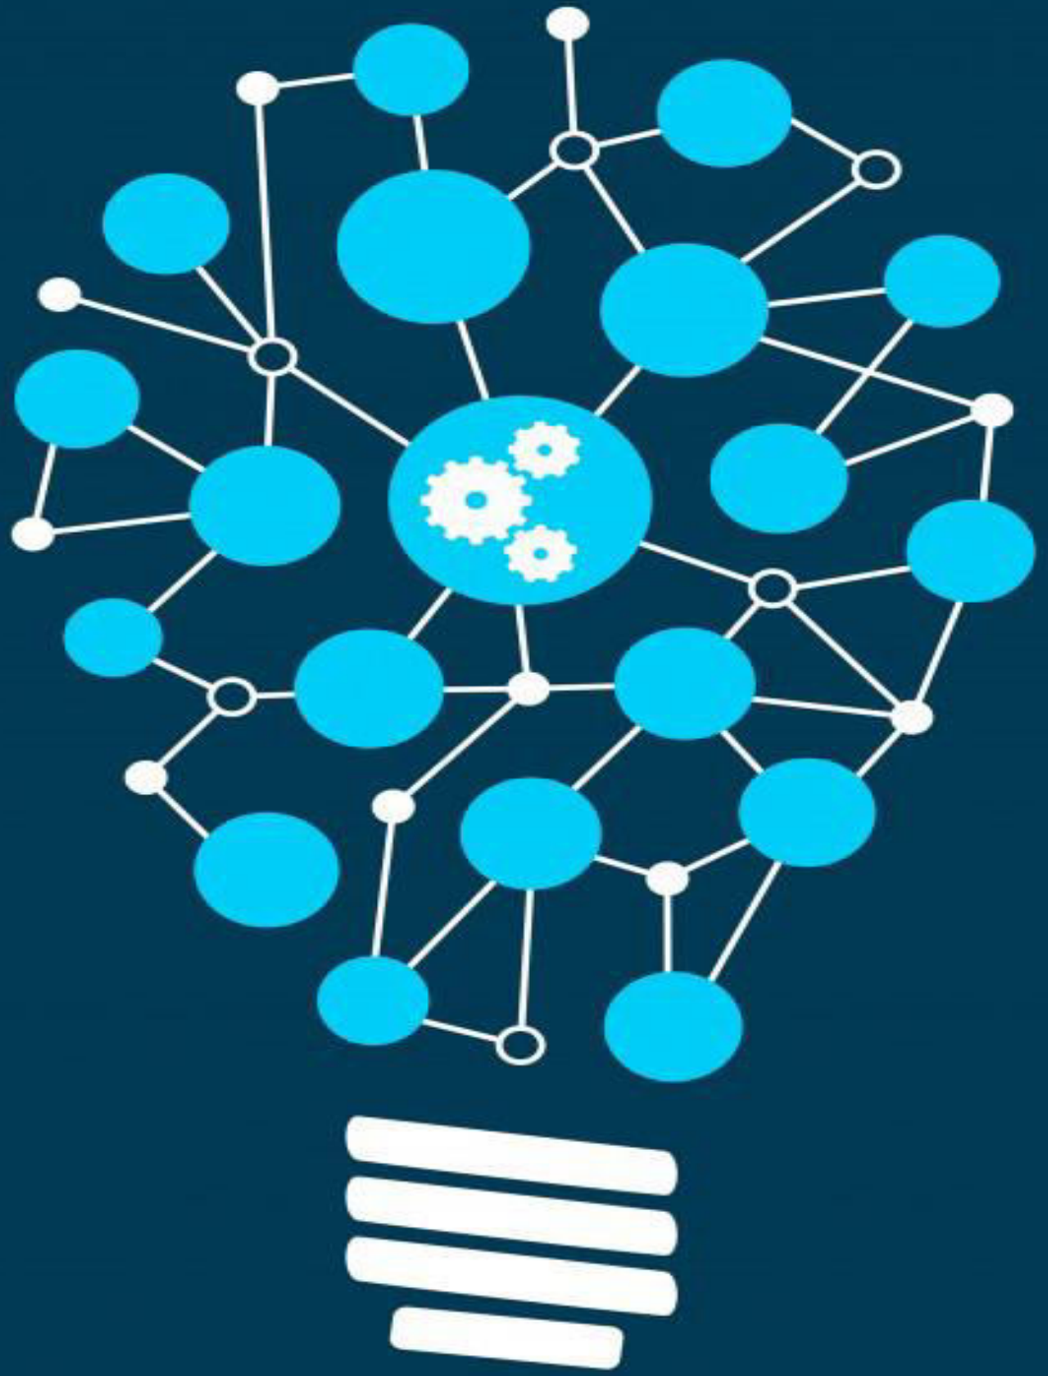

Microsoft

CRISP-DM

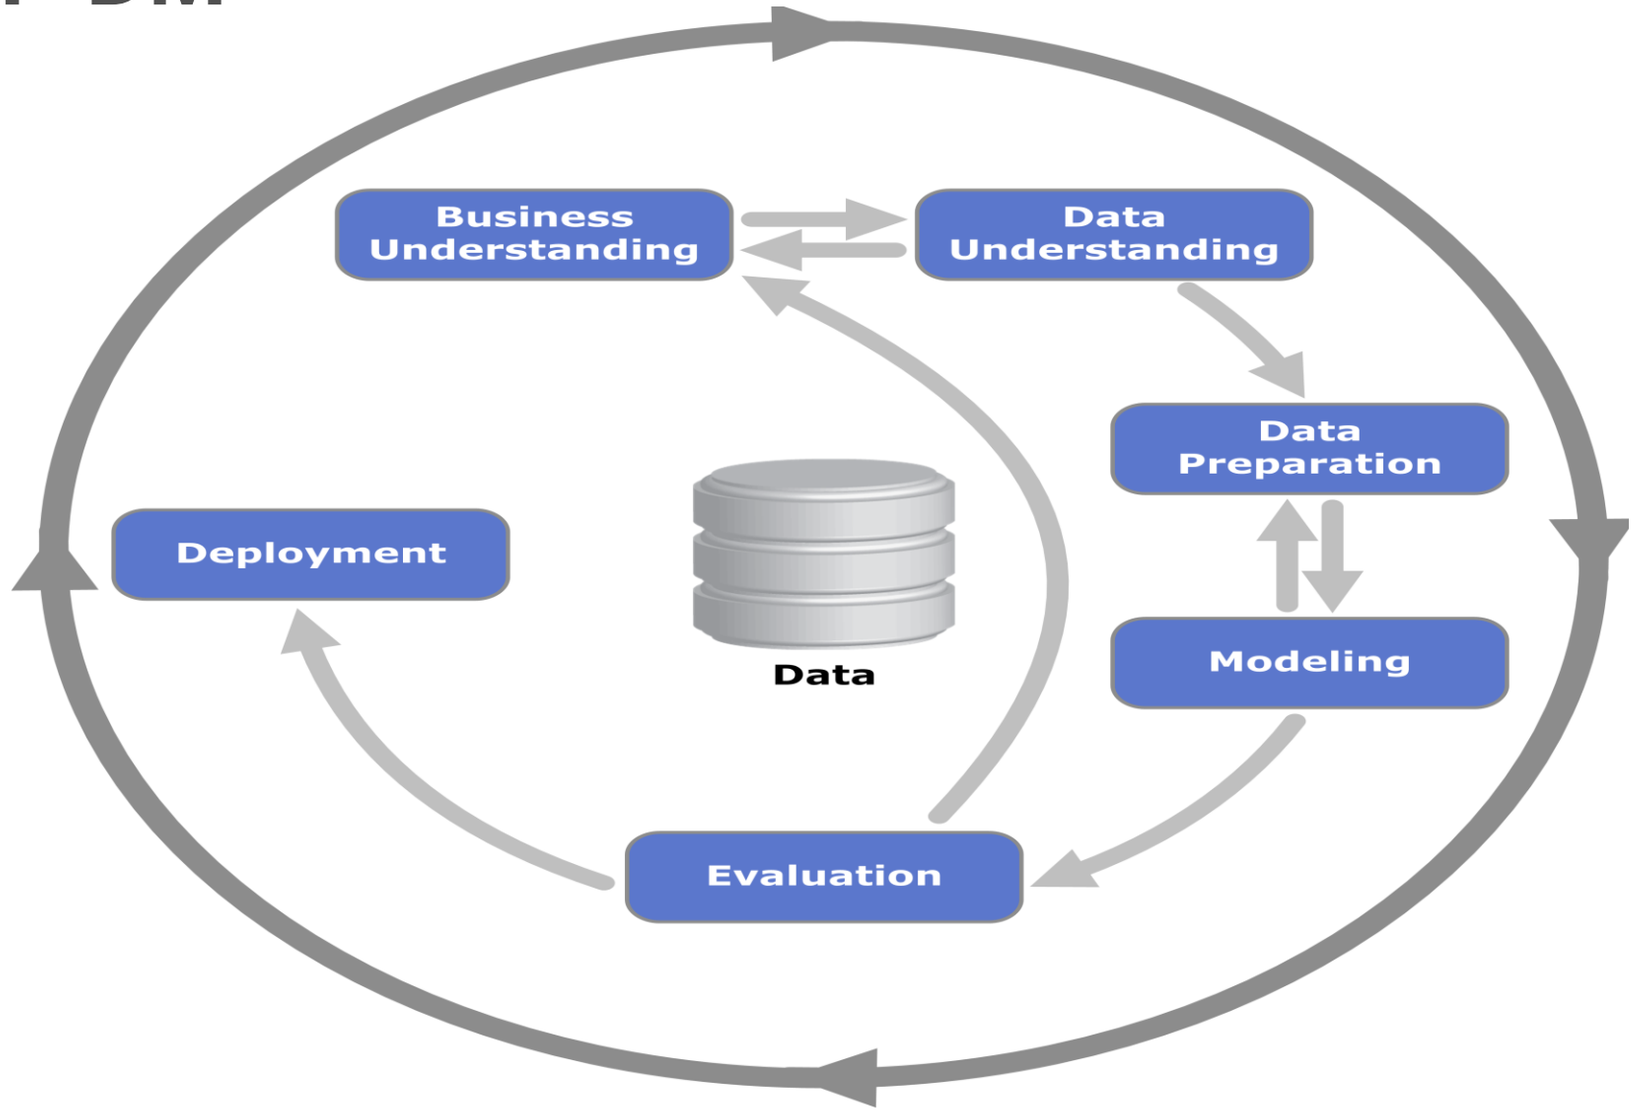

THE TEAM DATA SCIENCE PROCESS

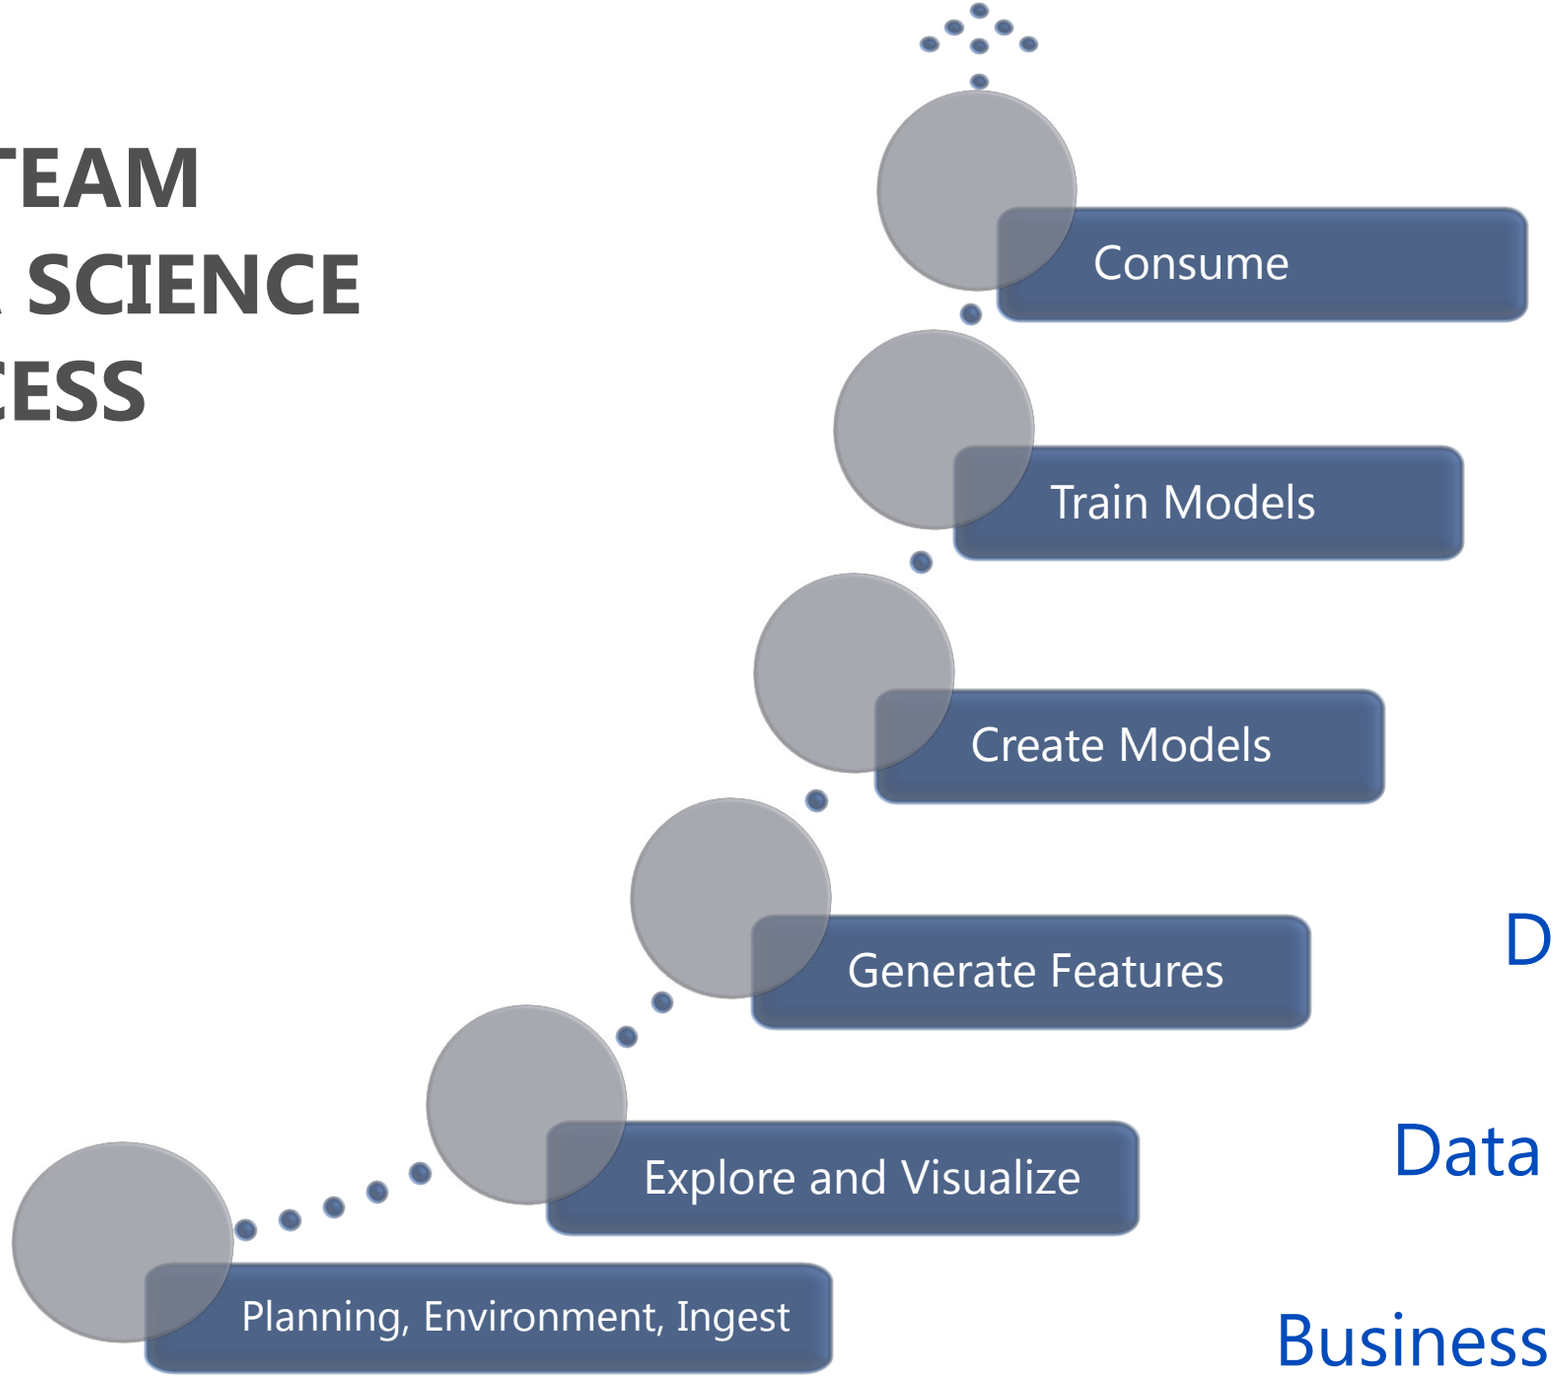

Deploying

Evaluating

Modeling

Data Preparation

Data Understanding

Business Understanding

Microsoft

CREATING AN EXPERIMENT

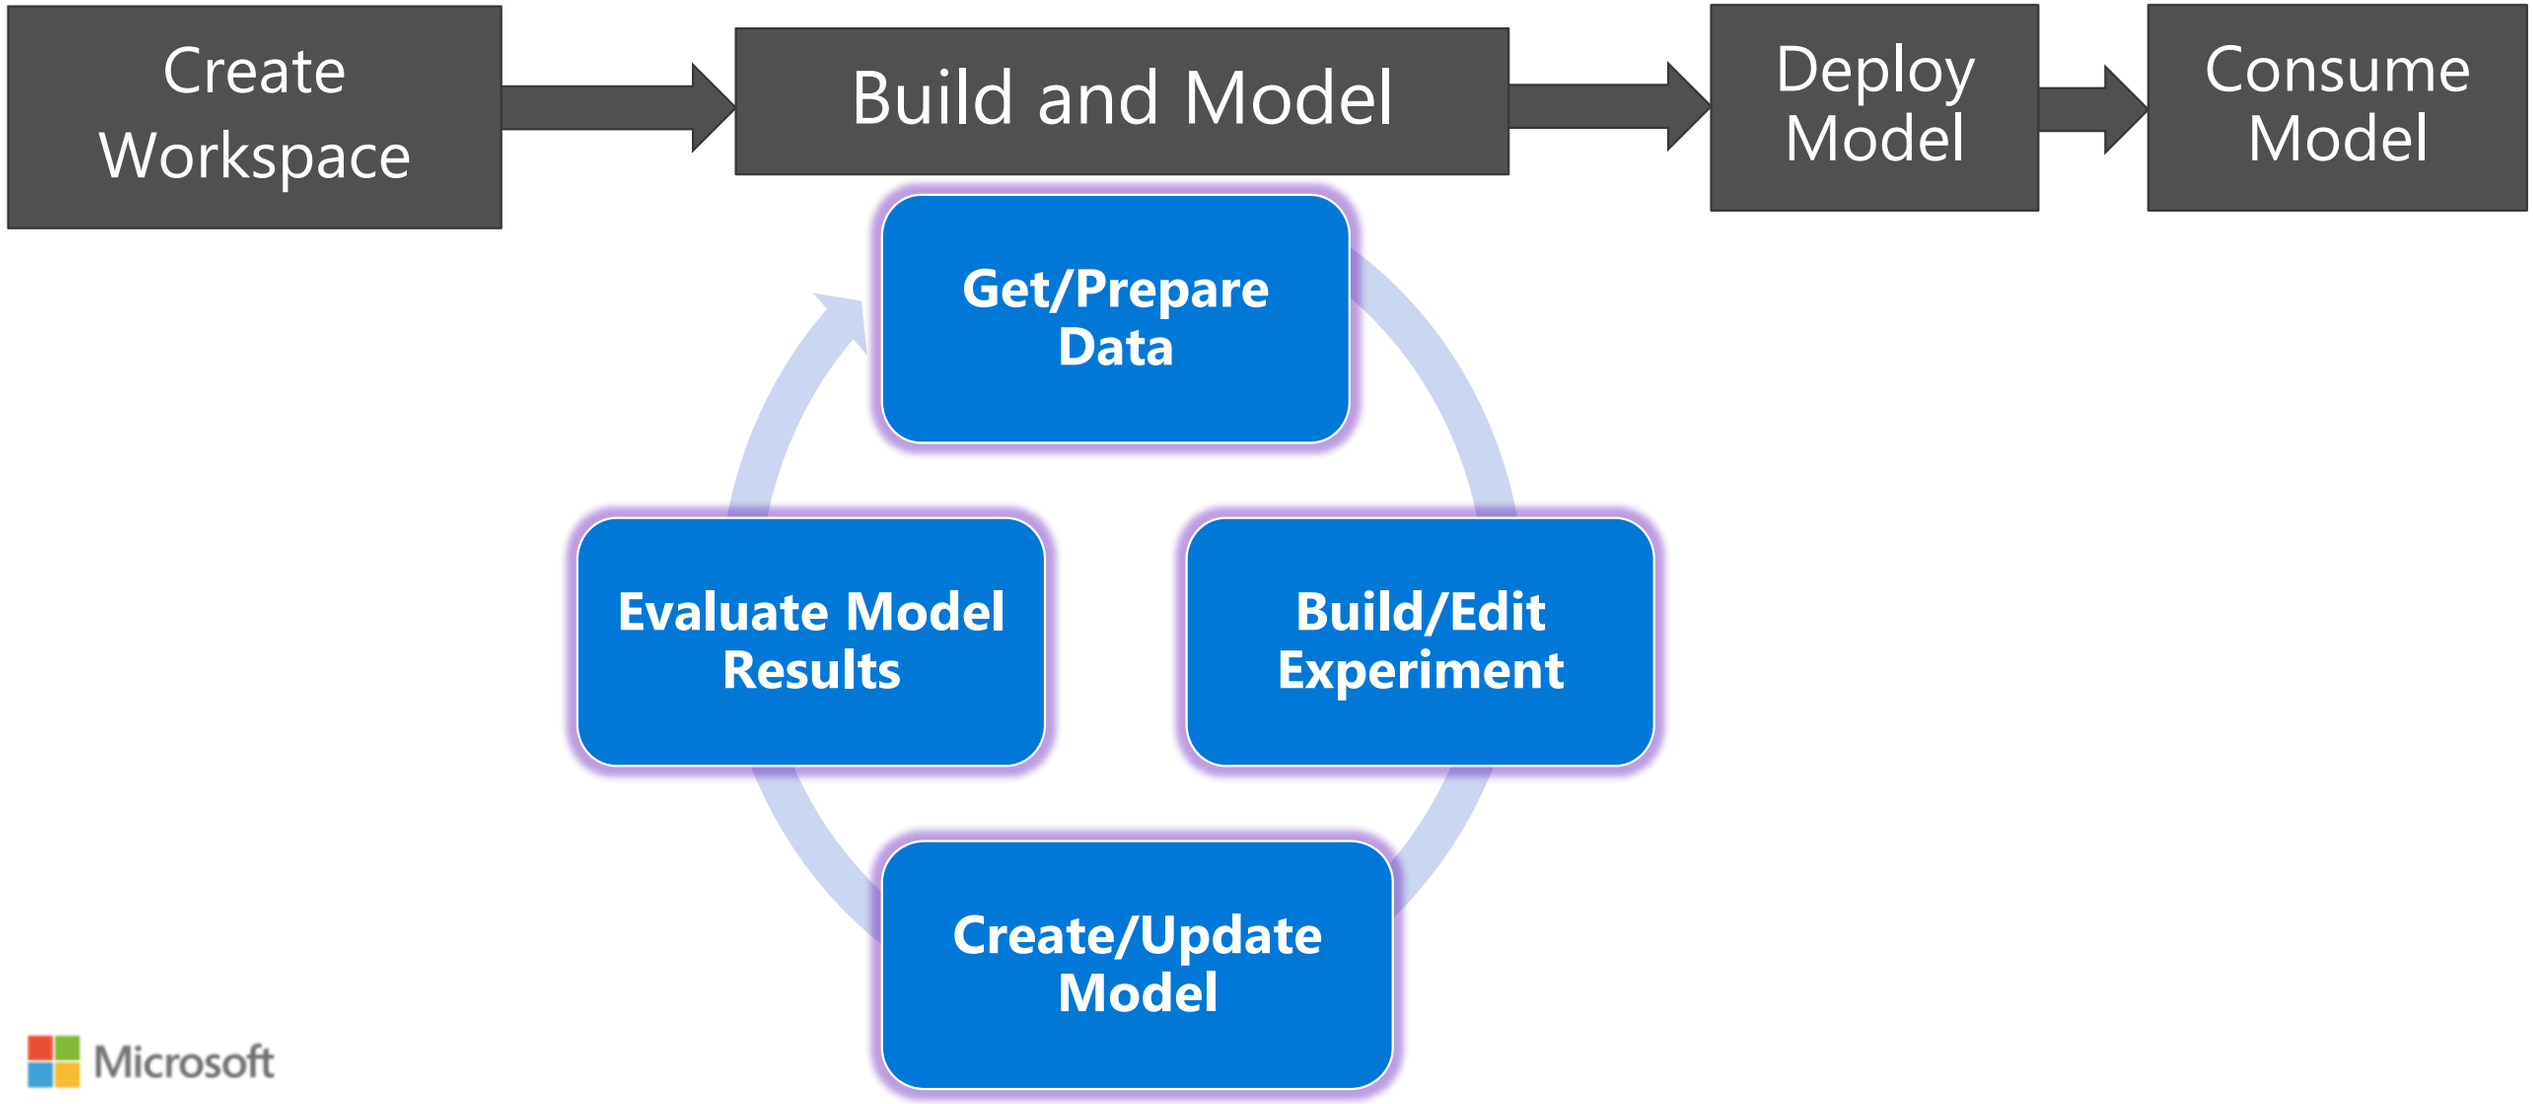

DATA INGESTION AND PREPARATION

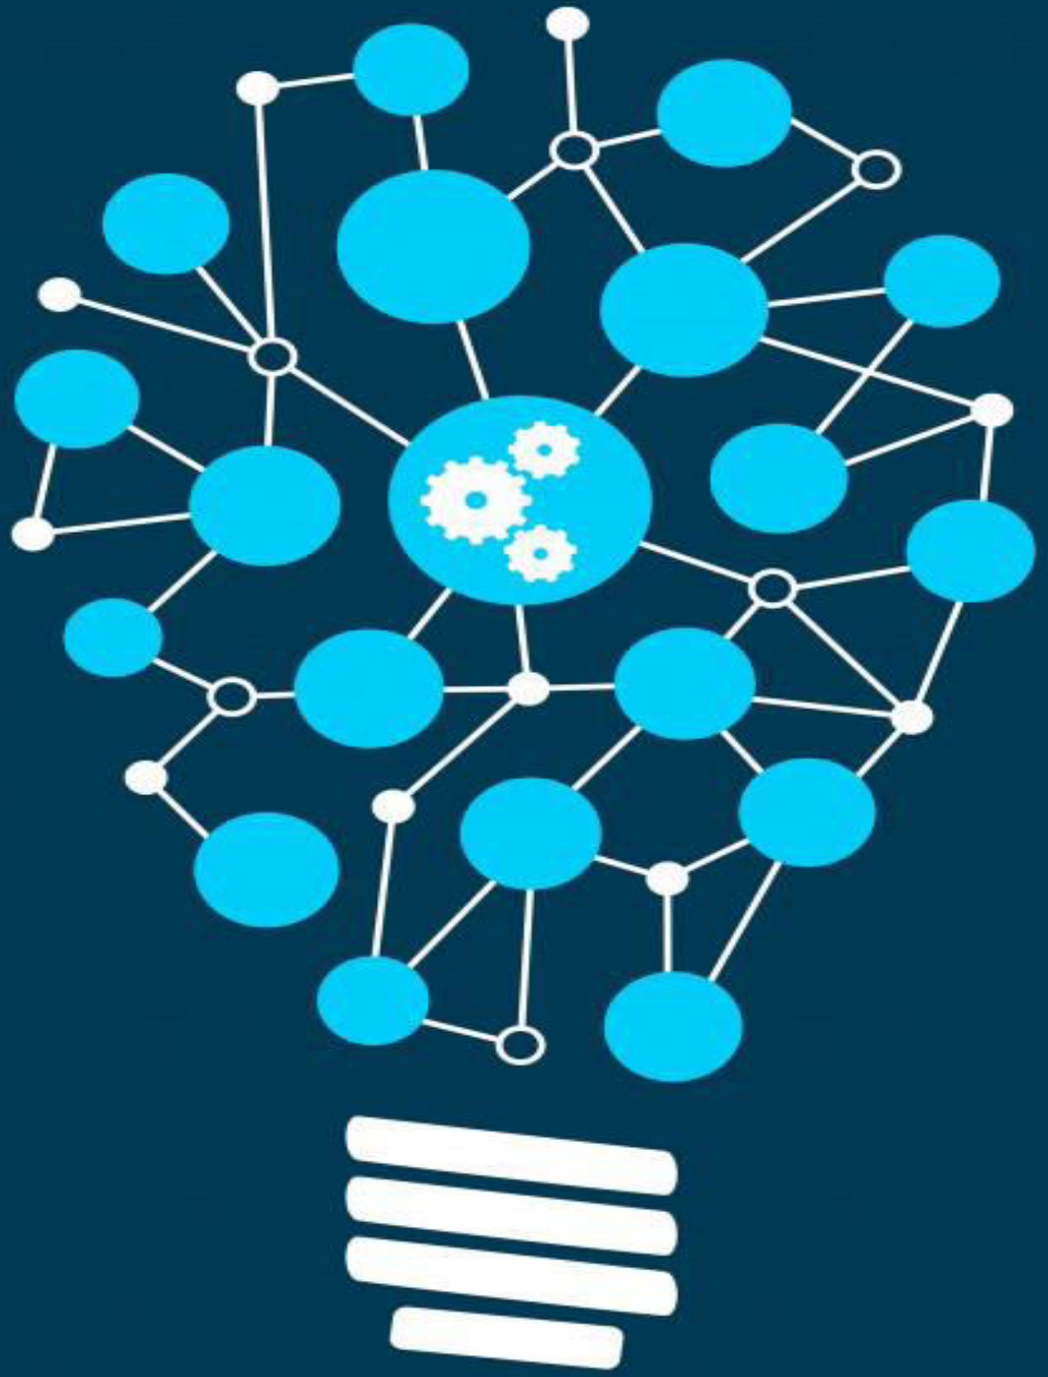

Microsoft

DATA ACCESS (IMPORT)

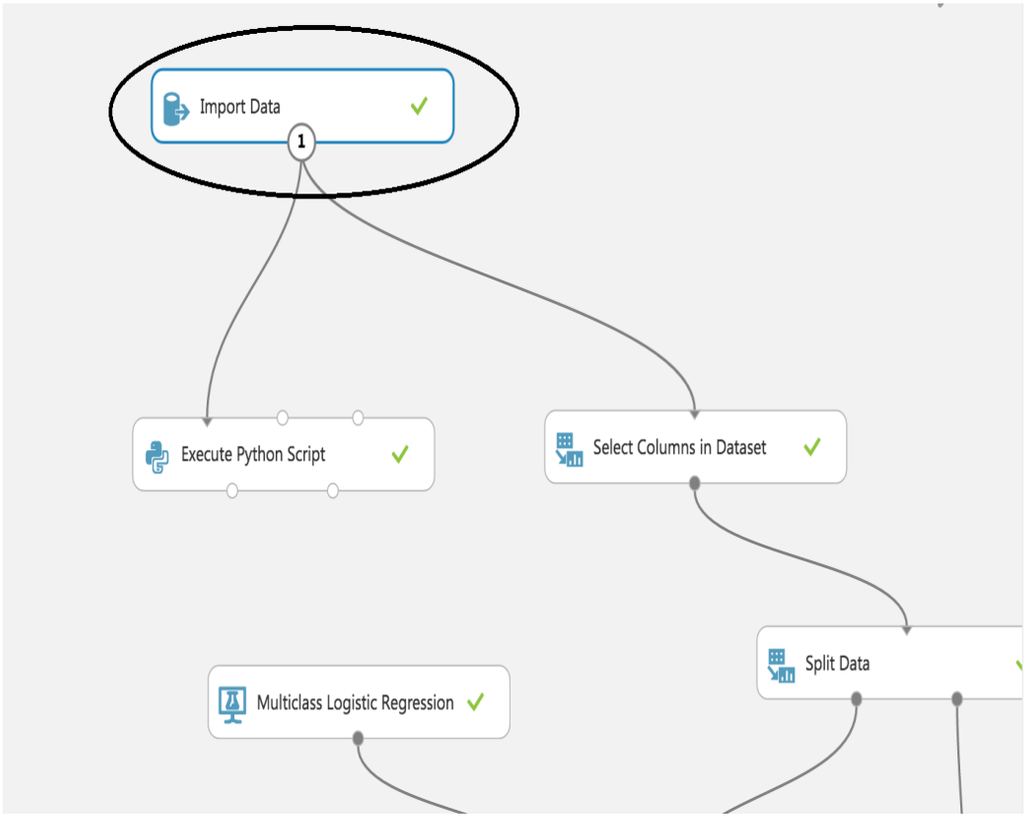

Import Data

Data source

Azure SQL Database

Database server name

irismldb3.database.windows.net

Database name

irisMLDB

User name

miprasad@irismldb3

Password

••••••••••

Accept any server certificate (insecure)

Database query

```
1 Select * from iris
```

Microsoft

DATA ACCESS (EXPORT)

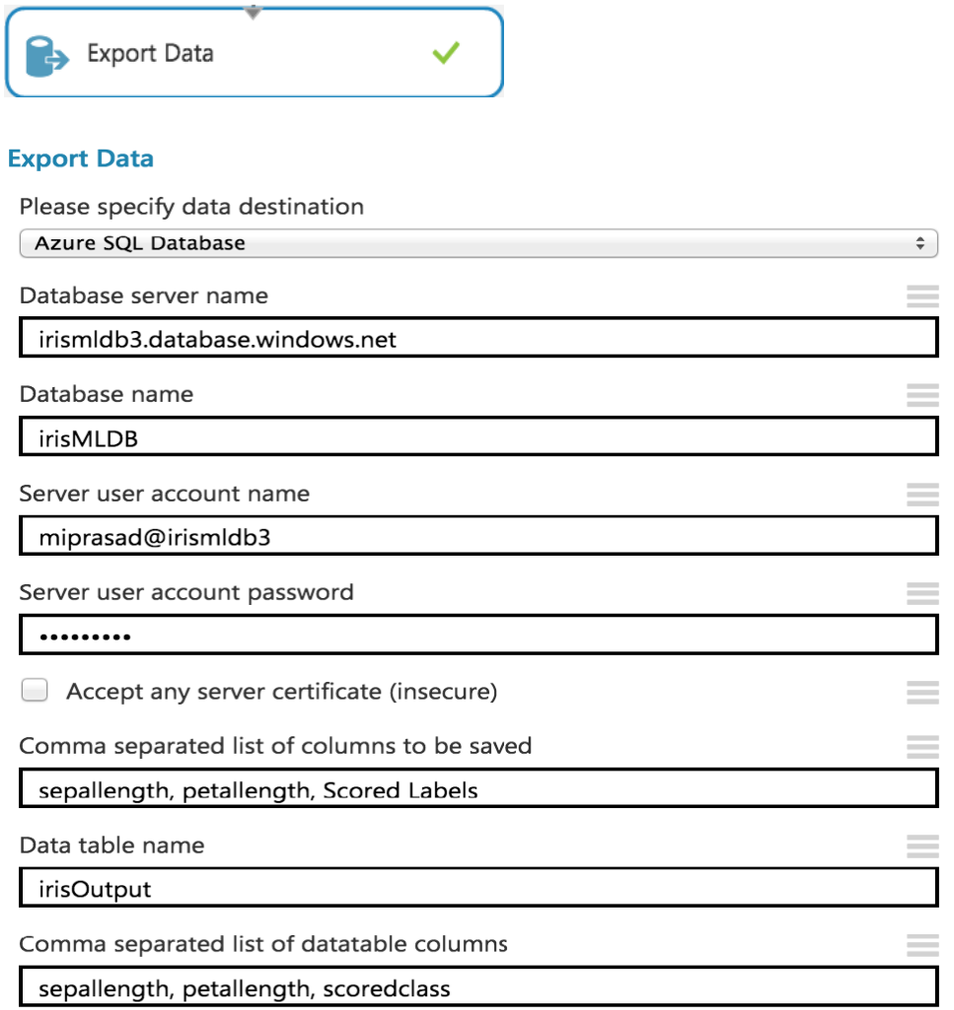

Data Format Conversion

- Convert to ARFF
- Convert to CSV
- Convert to Dataset
- Convert to SVMLight
- Convert to TSV

Microsoft

ALGORITHMS

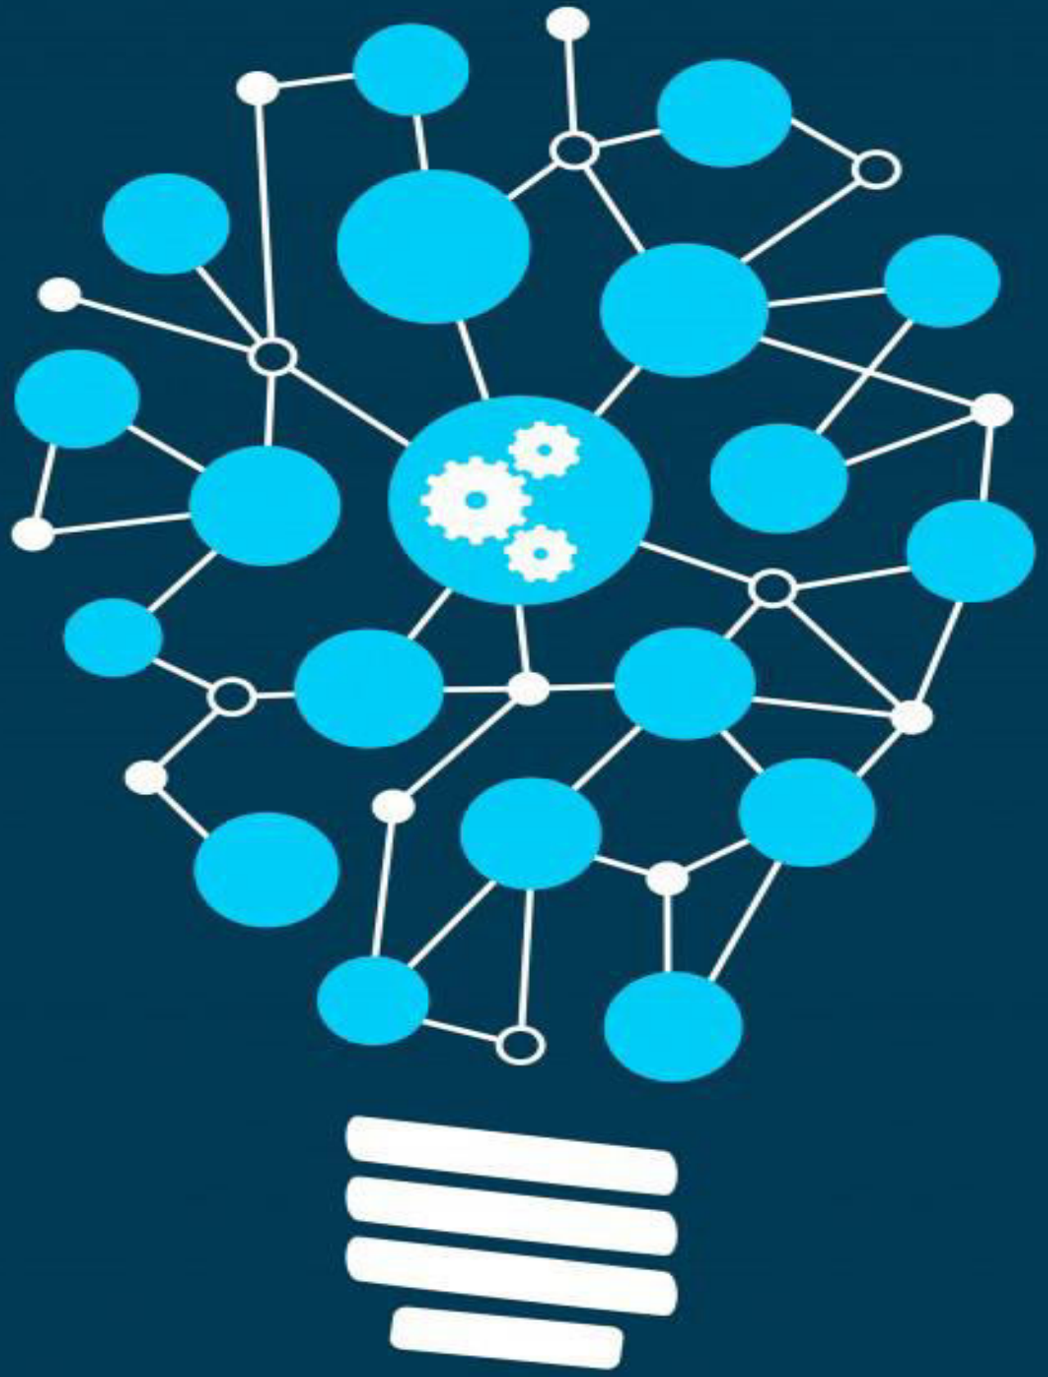

Microsoft

Microsoft Azure Machine Learning: Algorithm Cheat Sheet

This cheat sheet helps you choose the best Azure Machine Learning Studio algorithm for your predictive analytics solution. Your decision is driven by both the nature of your data and the question you're trying to answer.

START

Finding unusual data points

Discovering structure

Predicting categories

Predicting values

Three or more

Two

ANOMALY DETECTION

- **One-class SVM** — >100 features, aggressive boundary
- **PCA-based anomaly detection** — Fast training

CLUSTERING

- **K-means**

MULTI-CLASS CLASSIFICATION

- Fast training, linear model — **Multiclass logistic regression**
- Accuracy, long training times — **Multiclass neural network**
- Accuracy, fast training — **Multiclass decision forest**
- Accuracy, small memory footprint — **Multiclass decision jungle**
- Depends on the two-class classifier, see notes below — **One-v-all multiclass**

REGRESSION

- **Ordinal regression** — Data in rank ordered categories
- **Poisson regression** — Predicting event counts
- **Fast forest quantile regression** — Predicting a distribution
- **Linear regression** — Fast training, linear model
- **Bayesian linear regression** — Linear model, small data sets
- **Neural network regression** — Accuracy, long training time
- **Decision forest regression** — Accuracy, fast training
- **Boosted decision tree regression** — Accuracy, fast training, large memory footprint

TWO-CLASS CLASSIFICATION

- **Two-class SVM** — >100 features, linear model
- **Two-class averaged perceptron** — Fast training, linear model
- **Two-class logistic regression** — Fast training, linear model
- **Two-class Bayes point machine** — Fast training, linear model

- Accuracy, fast training — **Two-class decision forest**
- Accuracy, fast training, large memory footprint — **Two-class boosted decision tree**
- Accuracy, small memory footprint — **Two-class decision jungle**
- >100 features — **Two-class locally deep SVM**
- Accuracy, long training times — **Two-class neural network**

Microsoft

CLUSTERING

Grouping items based on defined Features

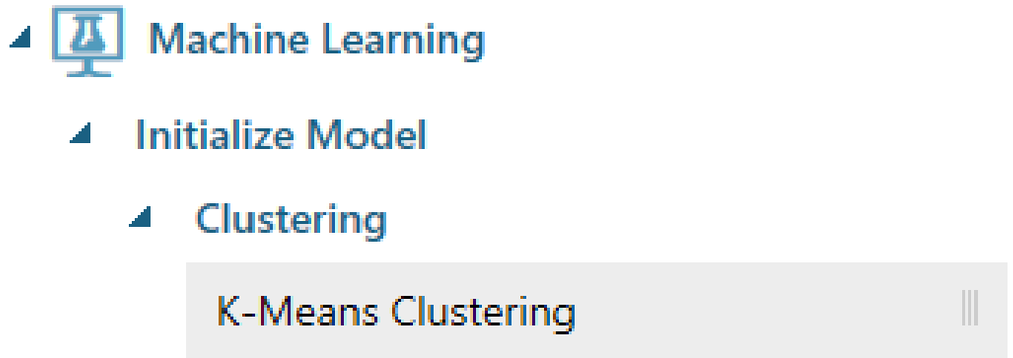

Microsoft

CLASSIFICATION

Predicting the class or category for a single instance of data

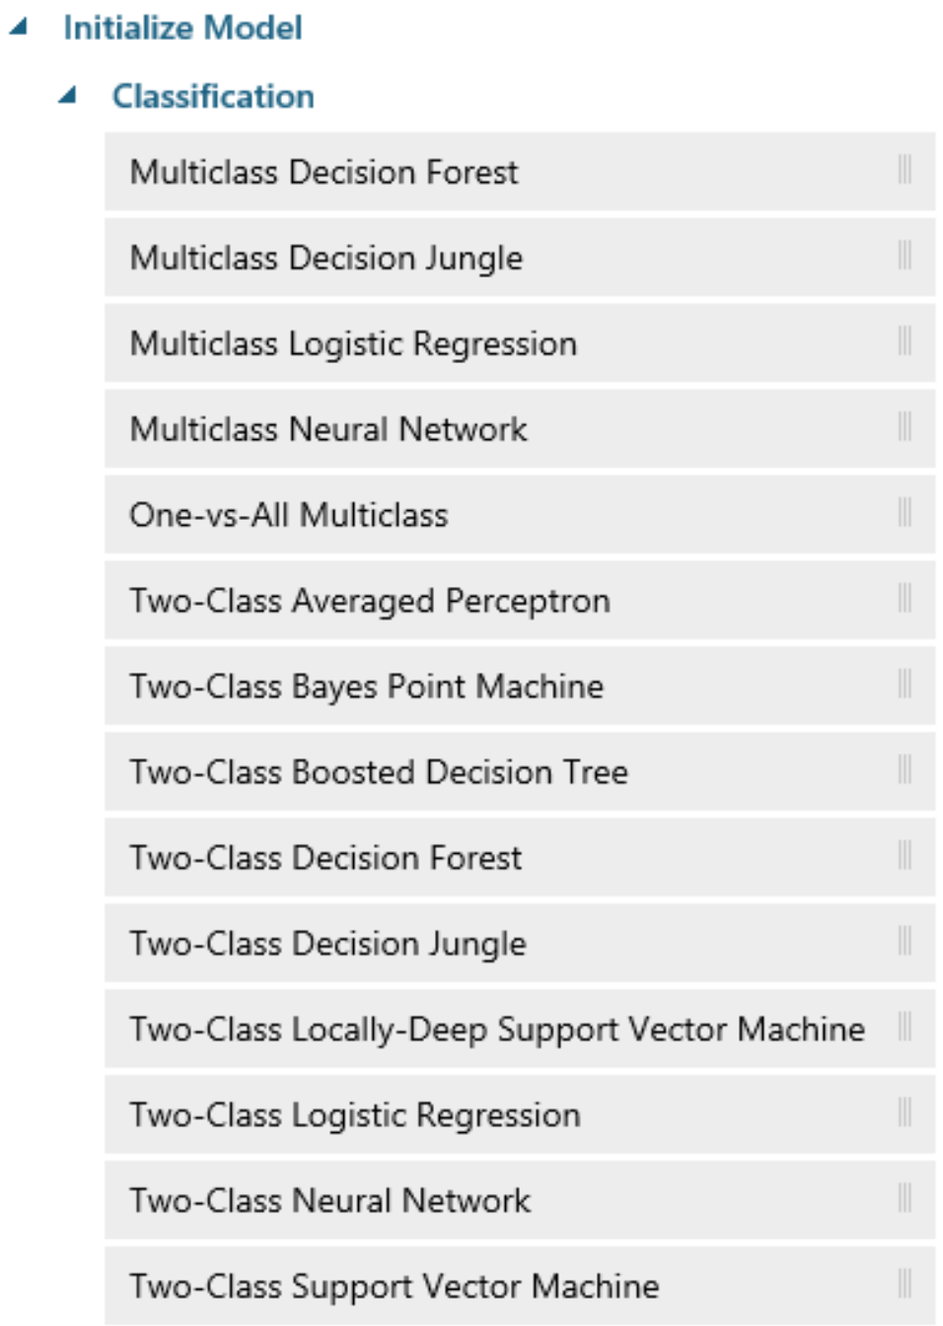

ANOMAY DETECTION

Selecting items based on unusual or suspicious patterns

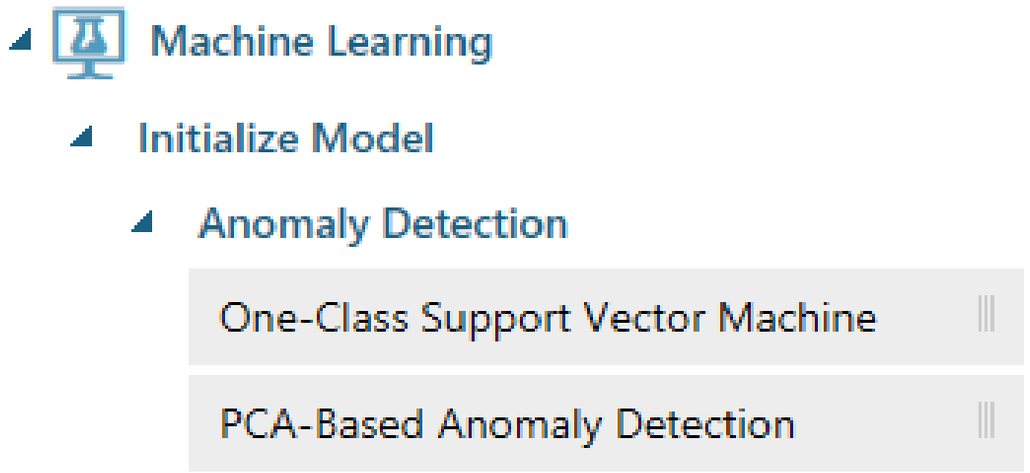

Microsoft

REGRESSION

Predicting the value of a datum given its history

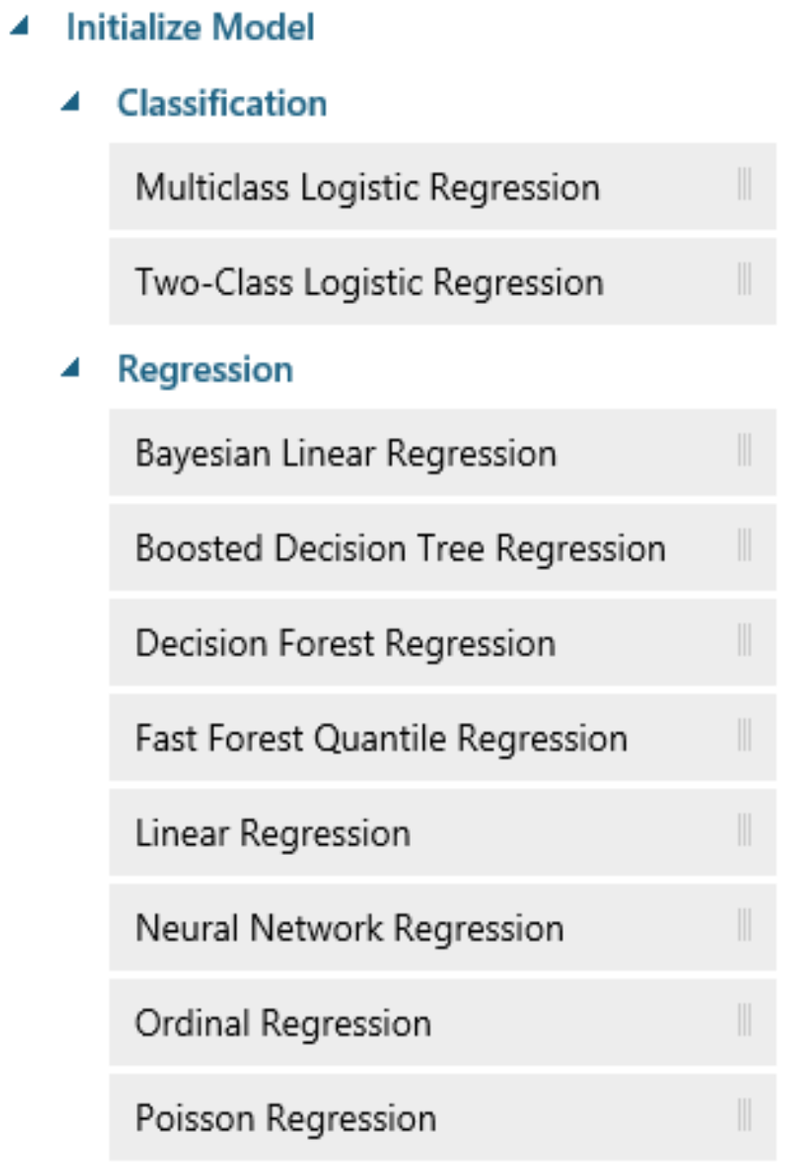

Microsoft

MODEL SCORING AND EVALUATION

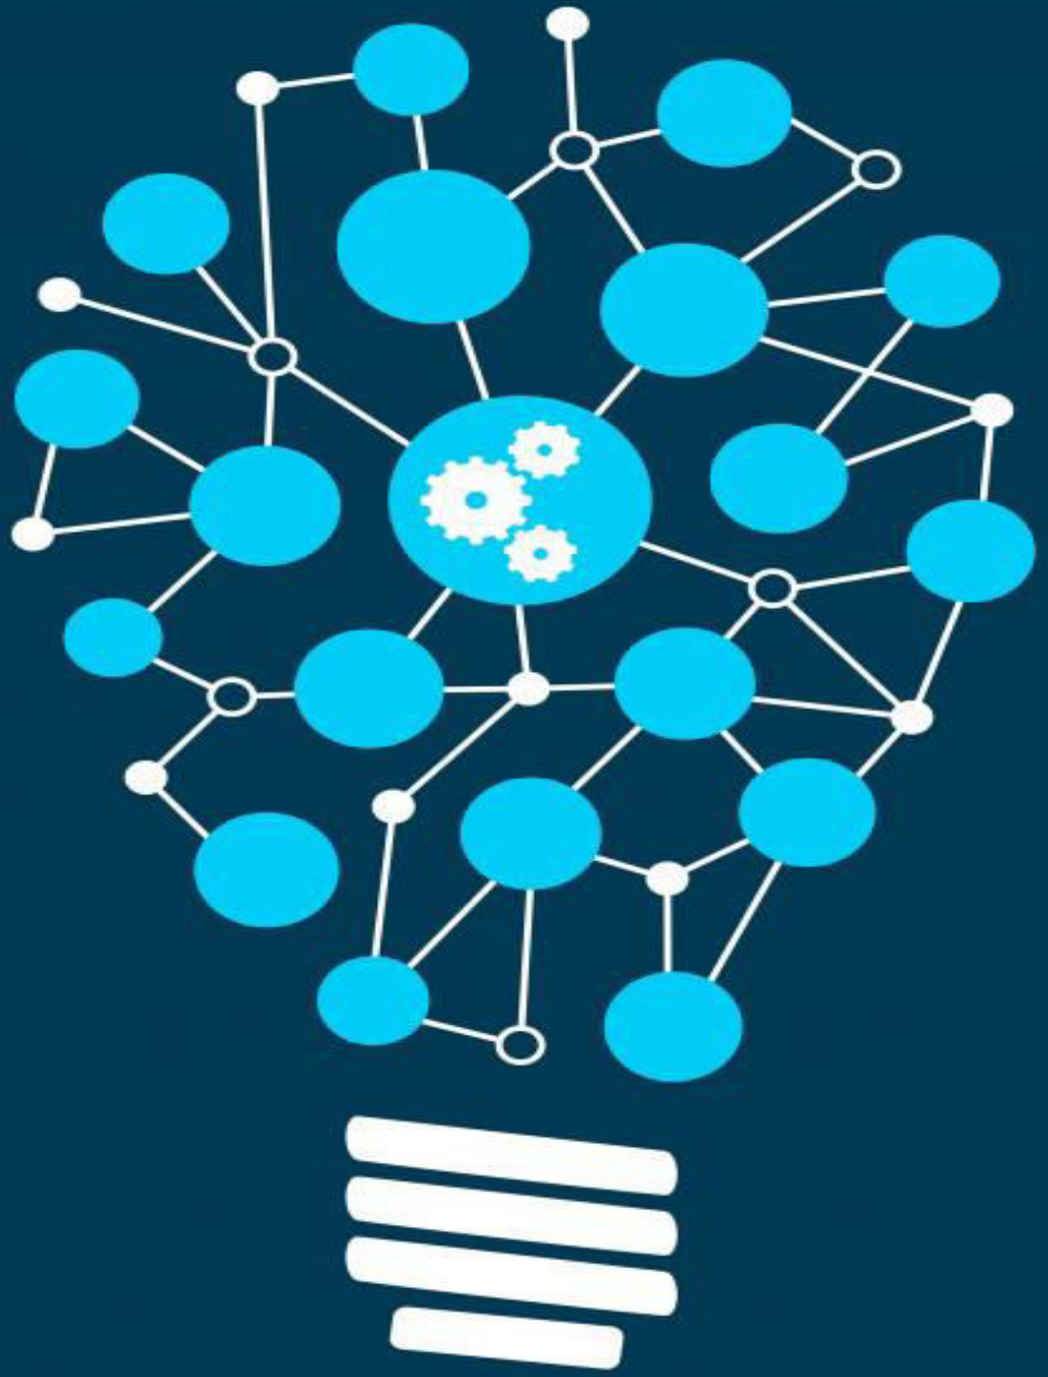

Microsoft

SCORING A MODEL

Apply a trained model to:

- A list of recommended items
- Forecasts for time series models
- Estimates of projected demand, volume, or other numeric quantity, for regression models
- Cluster assignments
- A predicted class or outcome, for classification models
- Probability scores associated with these outputs

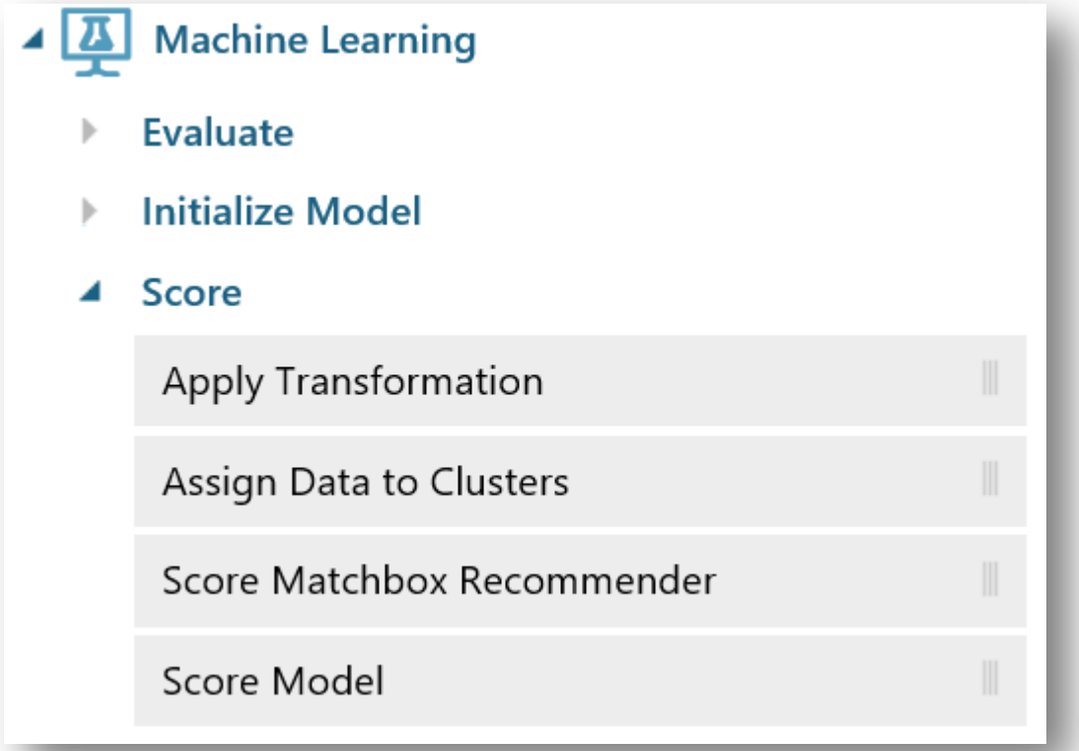

Microsoft

EVALUATING A MODEL

Metrics for Classification Models

- Accuracy, Recall, Precision, F1-Score
- AUC
- Average Log Loss
- Training Log Loss

Metrics for Regression Models

- Mean absolute error (MAE)
- Root mean squared error (RMSE)
- Relative absolute error (RAE)
- Relative squared error (RSE)
- Coefficient of determination

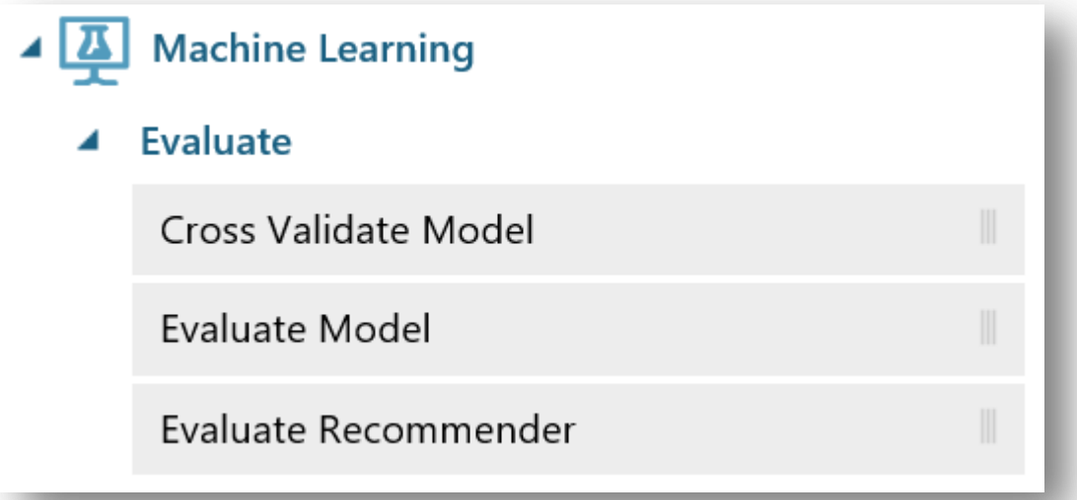

Microsoft

CROSS VALIDATION

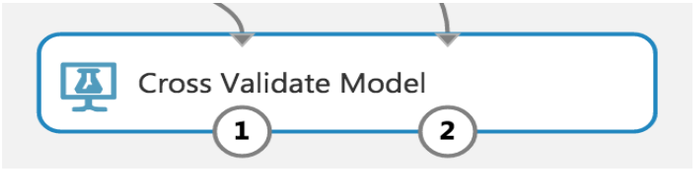

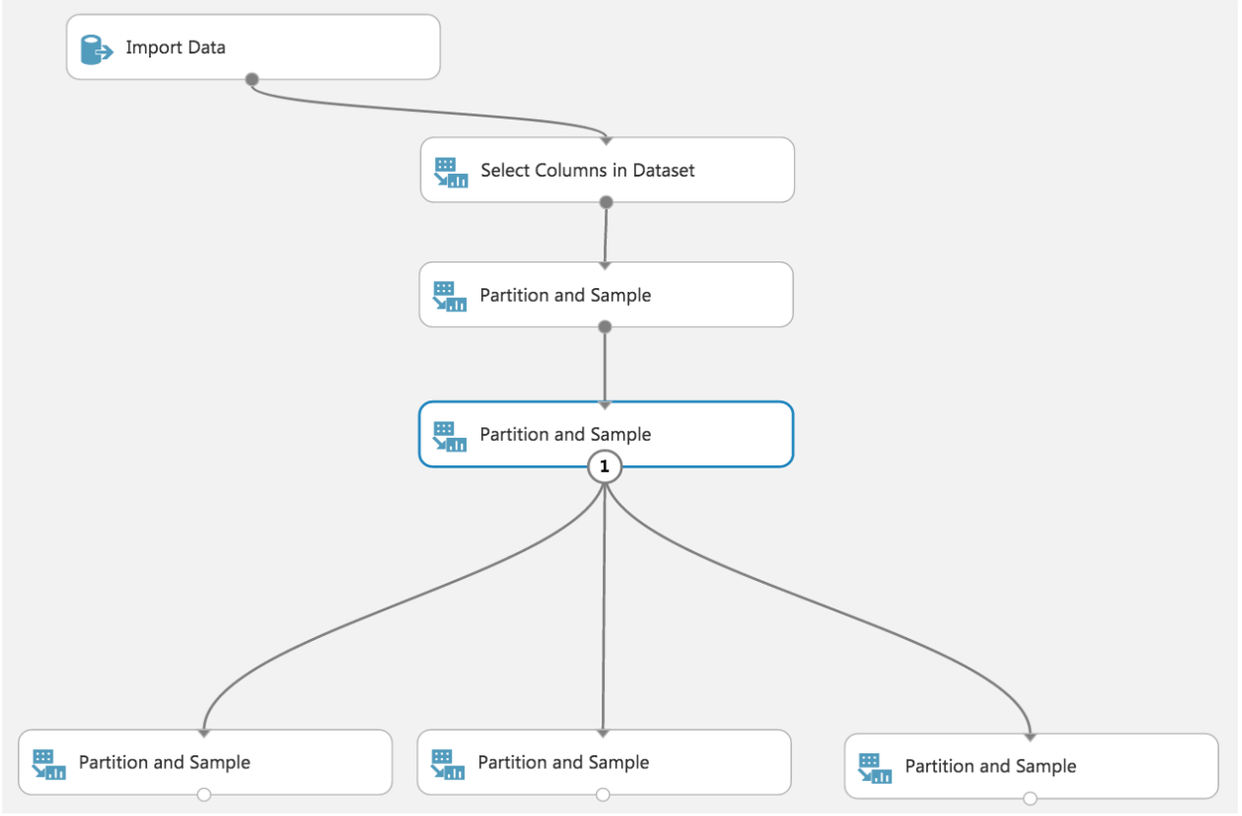

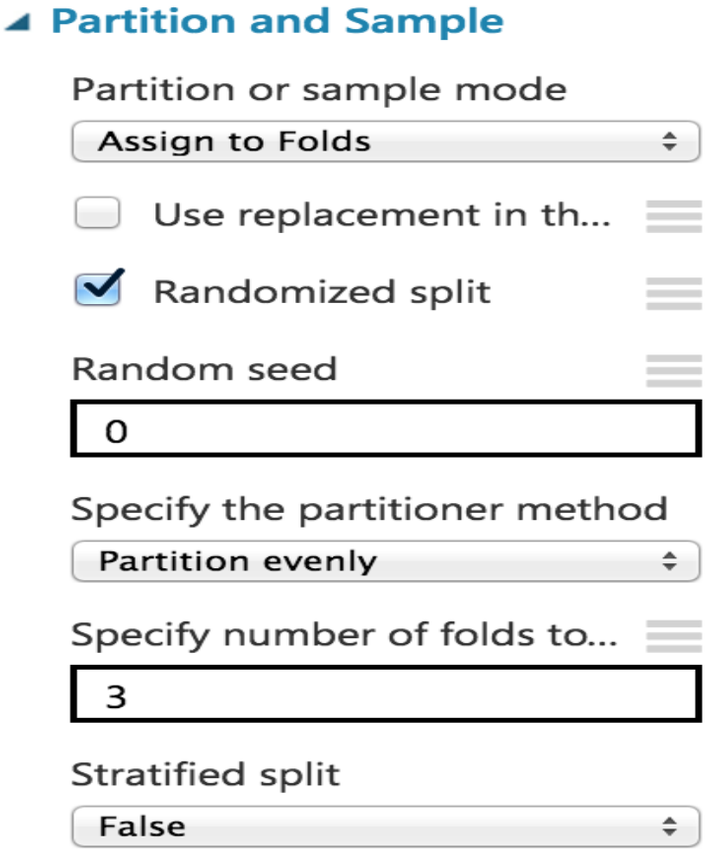

Microsoft

CUSTOMIZATION

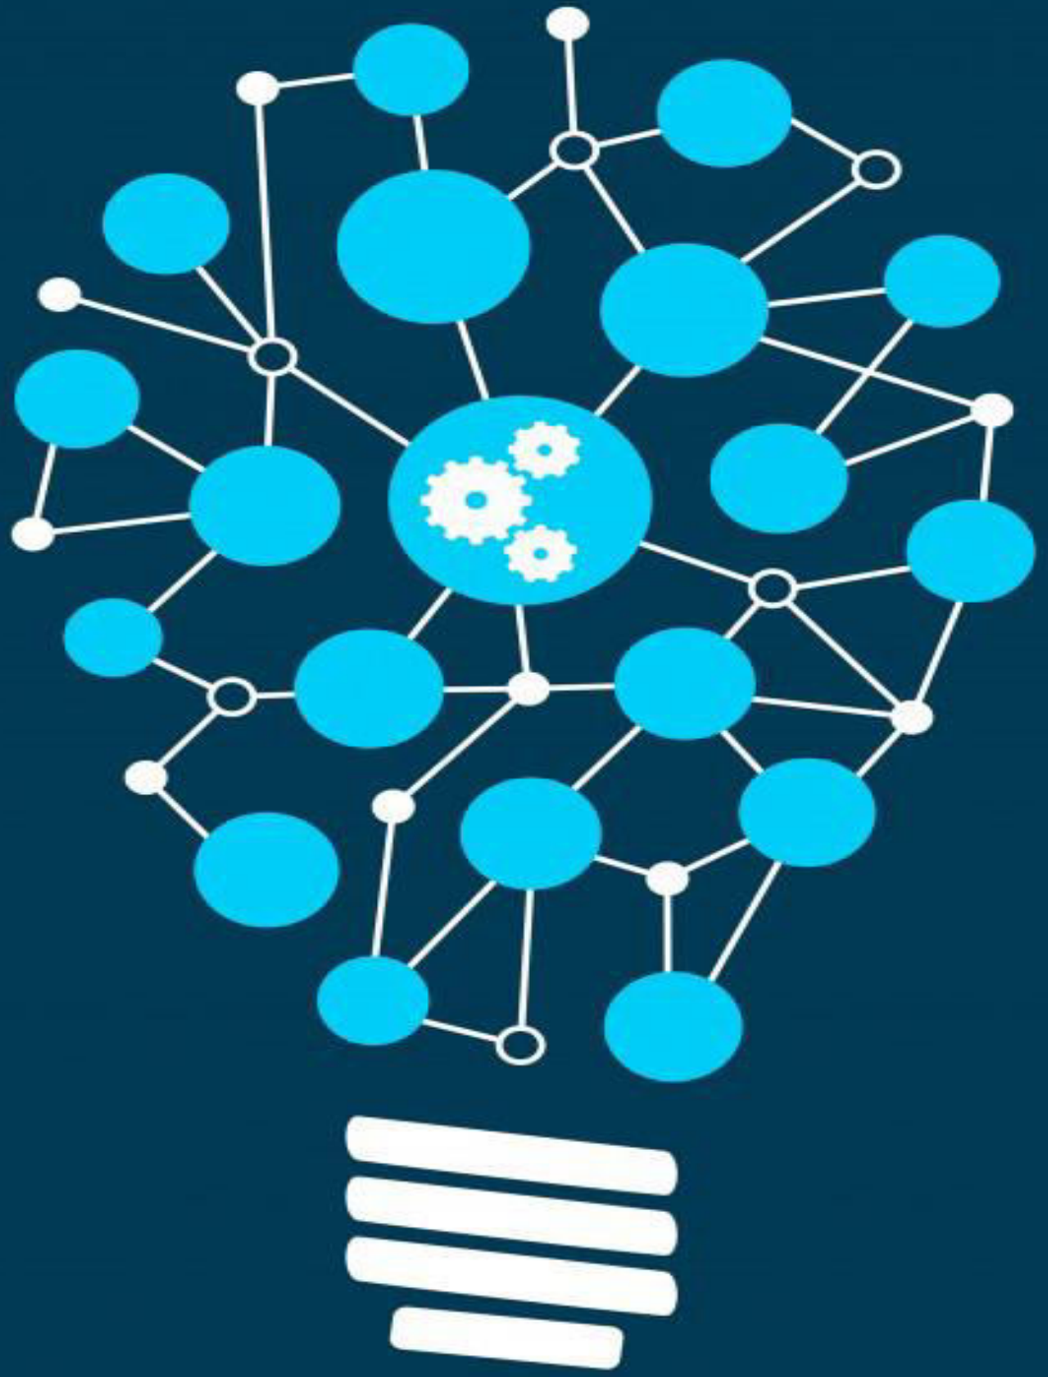

Microsoft

CUSTOM TASKS

Integrate Ipython notebooks with Azure Machine Learning to perform custom tasks:

- Visualization
- Use Python client libraries to enumerate datasets and models in your workspace
- Read, load, and manipulate data

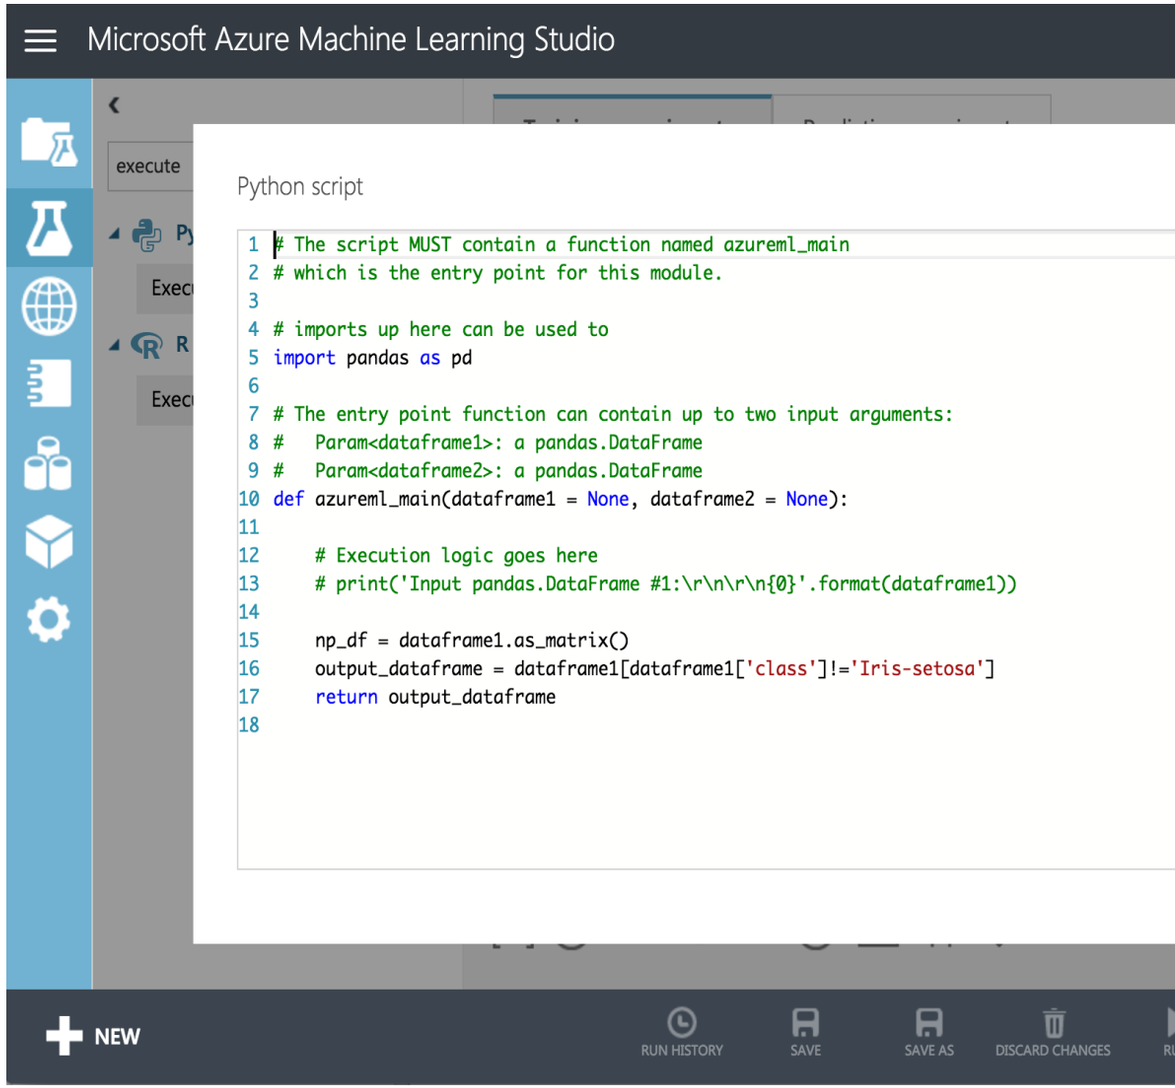

Microsoft

HOW TO USE EXECUTE PYTHON SCRIPT

1. Add the **Execute Python Script** module to your experiment.

2. Connect any datasets that you want to use for input. You can also provide a zipped file containing custom resources.

Dataset1. An optional dataset from your Machine Learning Studio workspace, containing input data or values.

Dataset2. A second dataset, also optional.

Script bundle. A zipped file containing custom resources.

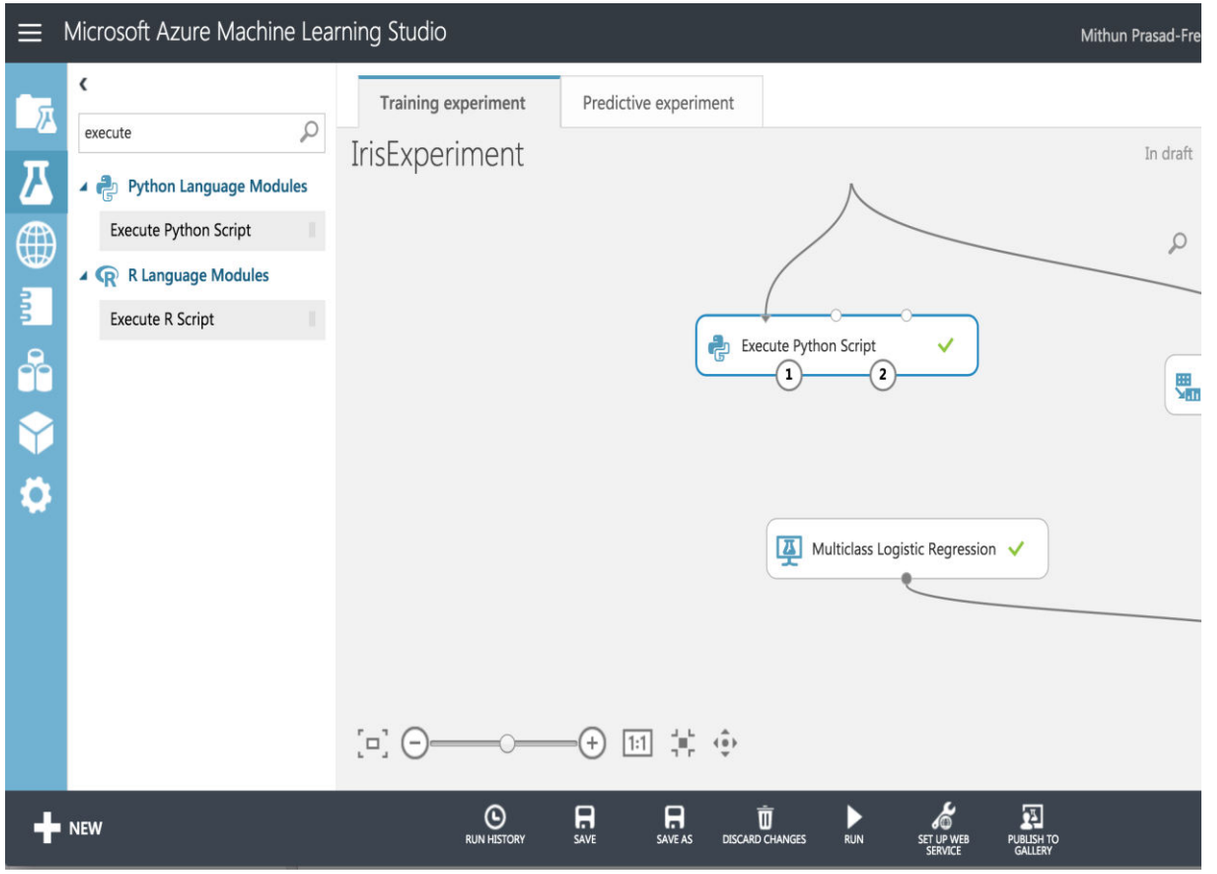

Microsoft